NSF/IUCRC CAC PROJECT

MONITORING, VISUALIZING, AND PREDICTING HEALTH STATUS OF HPC CENTERS

Jie Li
PhD Student, TTU
08/30/2019

Advisors:
Mr. Jon Hass, SW Architect, Dell Inc.
Dr. Alan Sill, Managing Director, HPCC, TTU
Dr. Yong Chen, Associate Professor, CS Dept, TTU
Dr. Tommy Dang, Assistant Professor, CS Dept, TTU

AGENDA

- Overview of the project
- Interleaving software metrics with hardware metrics

OVERVIEW

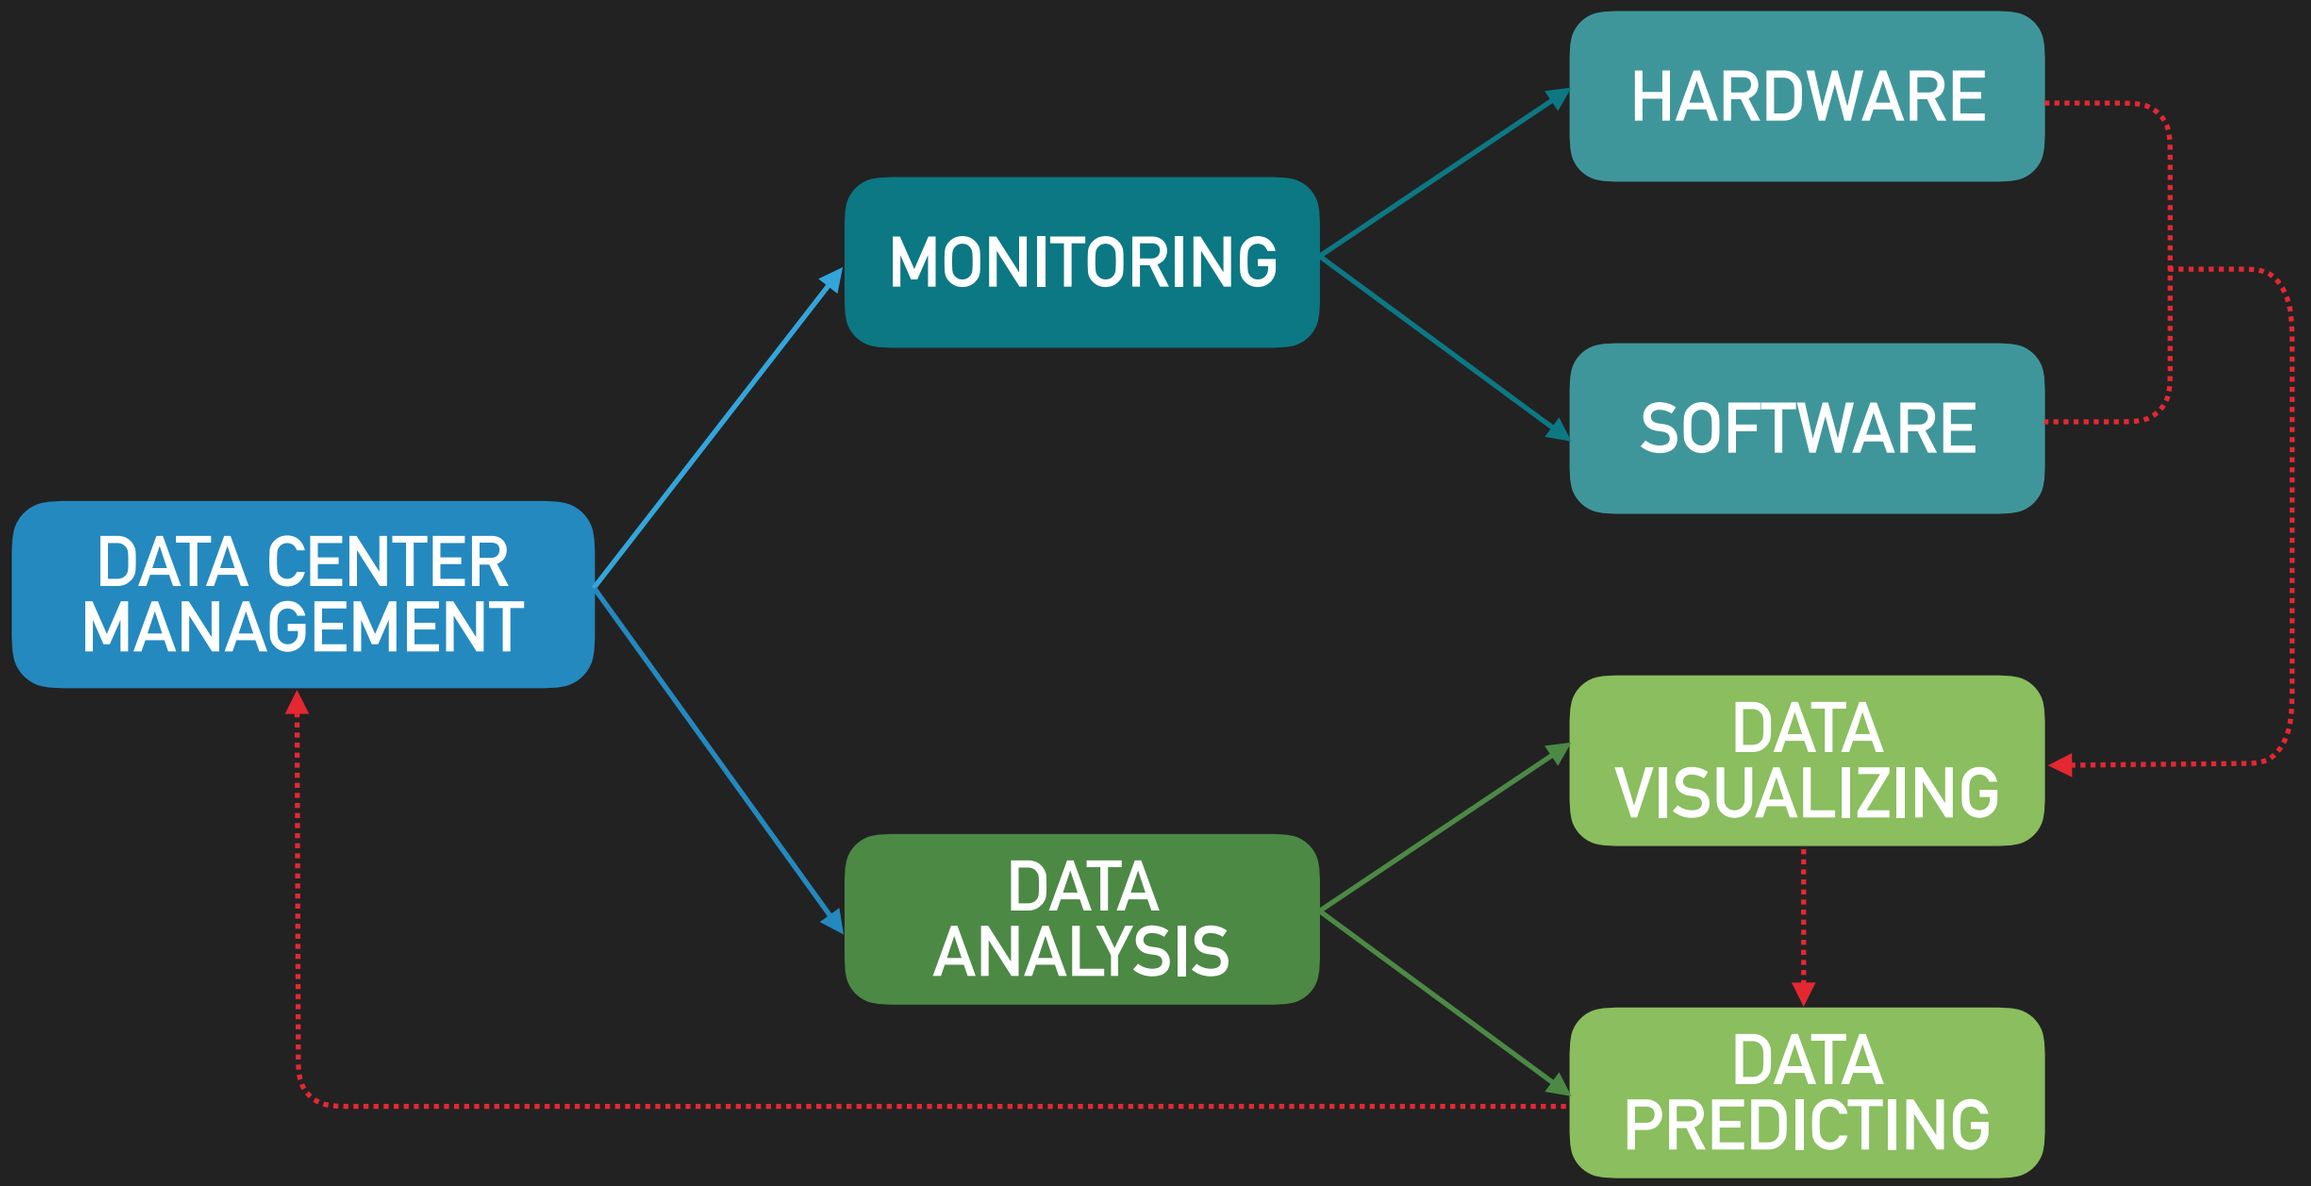
MONITORING
HARDWARE
SOFTWARE
DATA CENTER MANAGEMENT
DATA ANALYSIS
DATA VISUALIZING
DATA PREDICTING

MONITORING

- IMPROVE the ways in which systems and applications operate
- Measured in terms of performance, reliability, power usage, the ability to meet service level agreements, and similar metrics important to applications and IT infrastructure providers
- Operate across multiple time scales, across different size systems, and at multiple level of abstraction(application-centric, infrastructure-centric)
- Understand and analyze captured data, in addition to support intelligent problem determination methods, as needed by subsequent management actions. [1]

Ref: [1] Mahendra Kutare et al. Monalytics: Online Monitoring and Analytics for Managing Large Scale Data Centers. ICAC, 2010

MONITORING FRAMEWORK

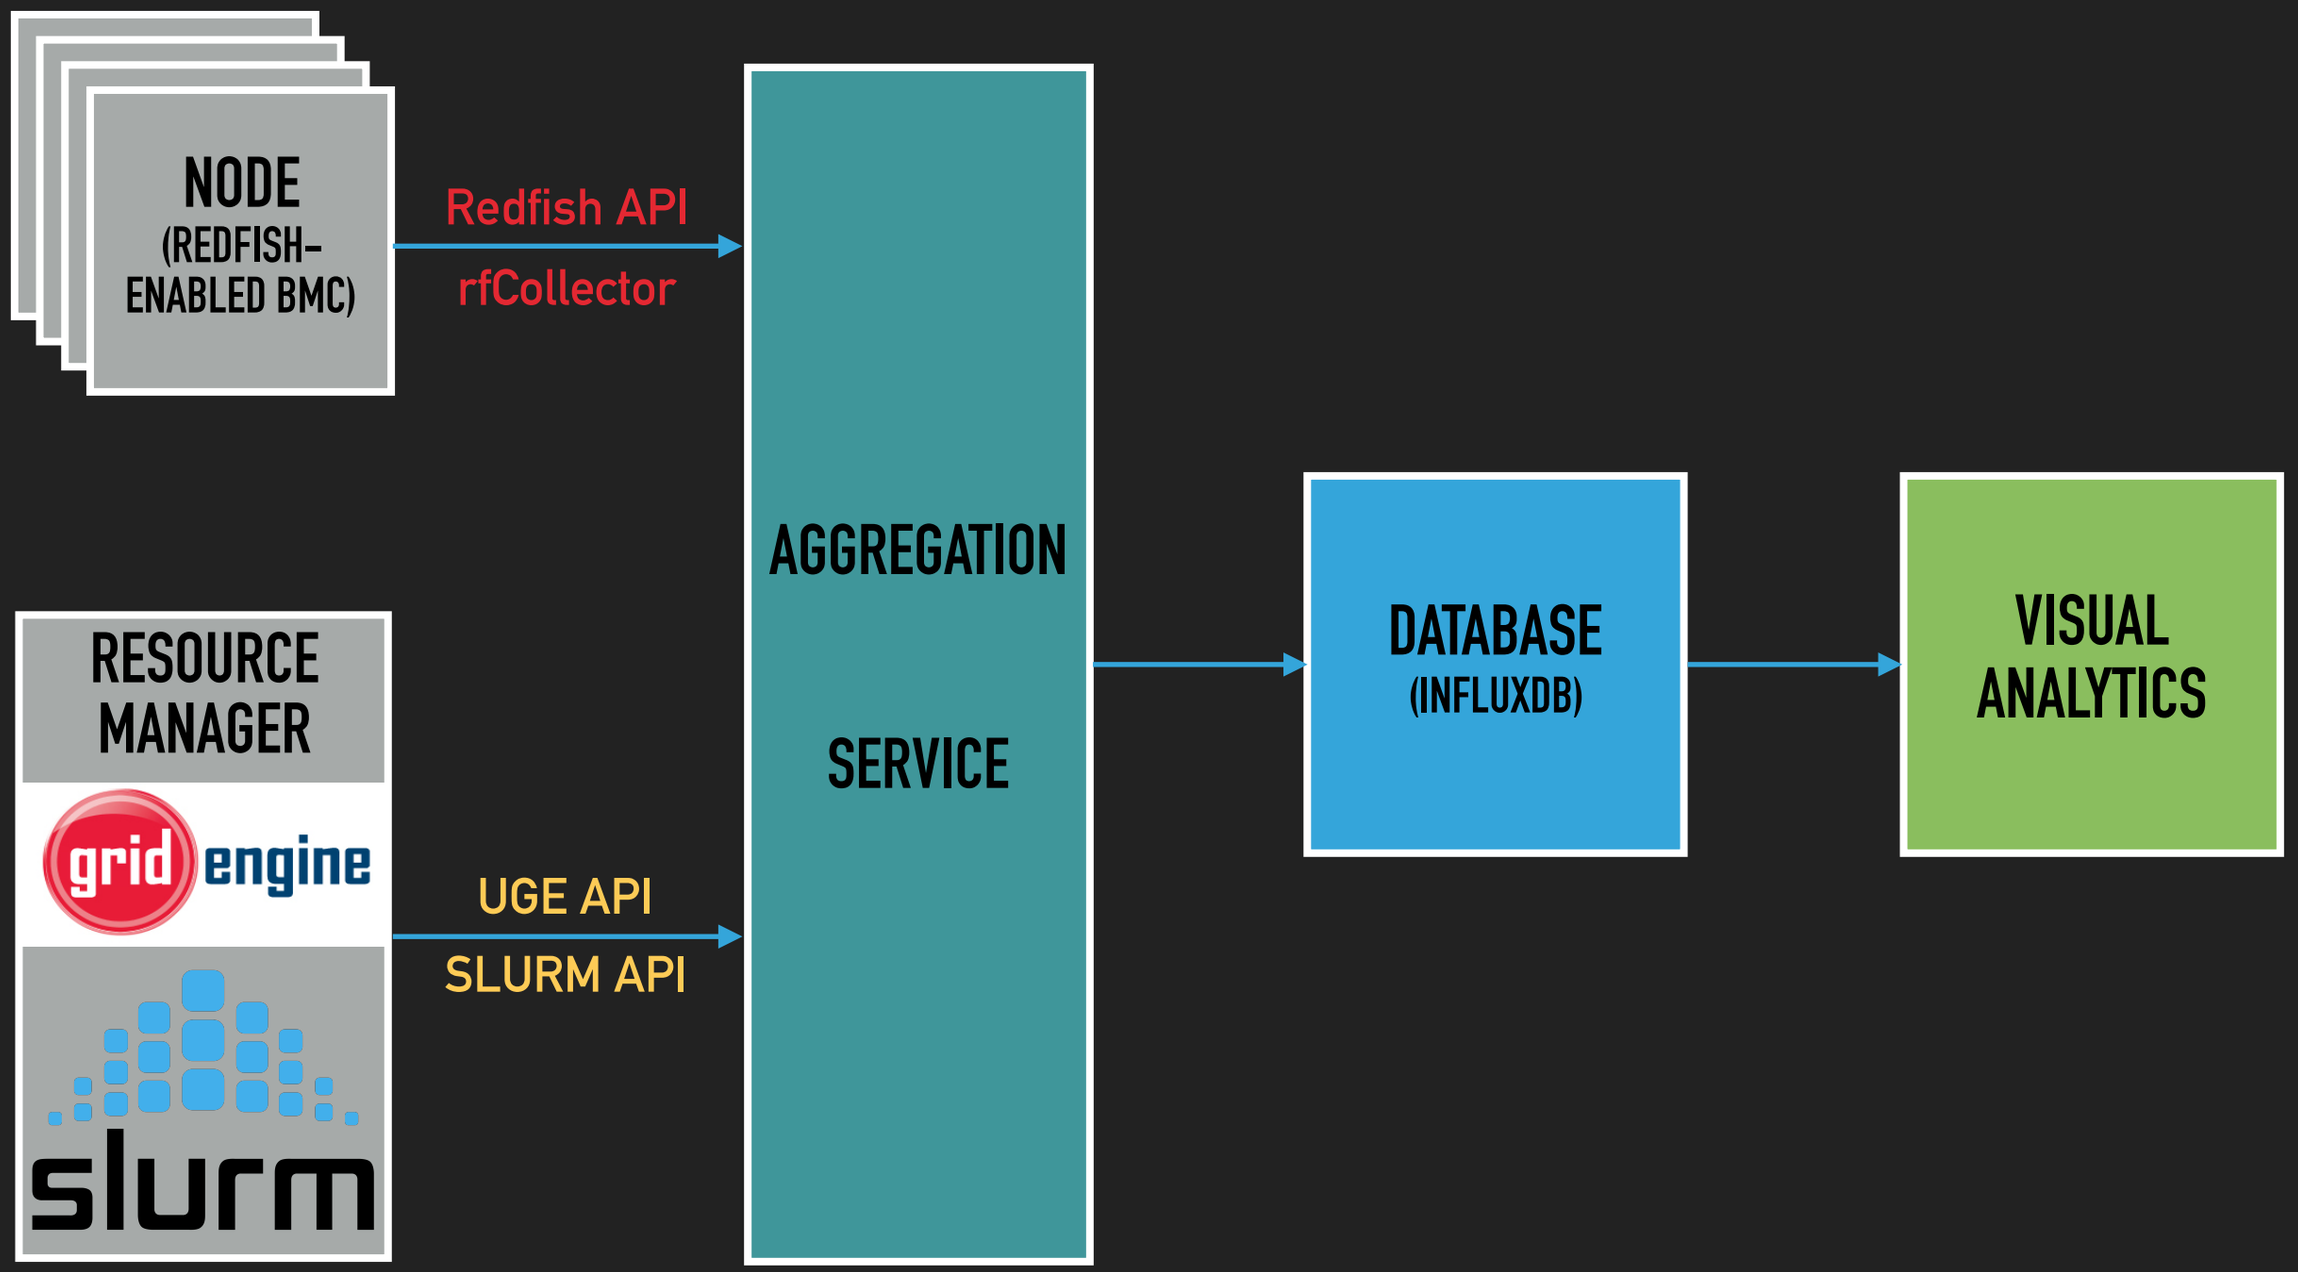

AGGREGATION SERVICE

AGGREGATION SERVICE

- Queries and Collects data across the entire cluster
- Builds metrics after receiving the monitoring data
- Interleaving BMC metrics with Resource Manager metrics
- Stores metrics in database(InfluxDB)

AGGREGATION SERVICE

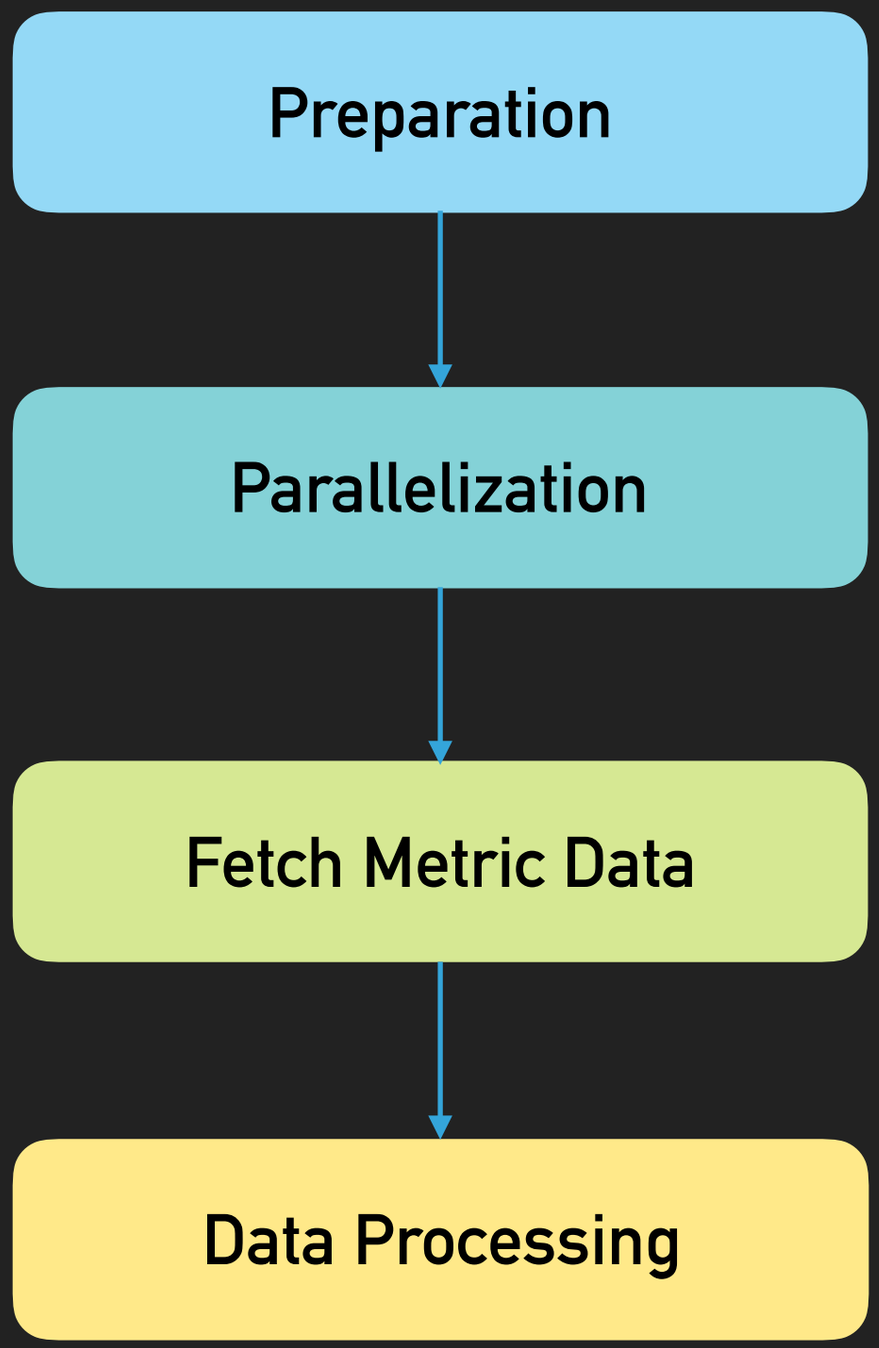

- Builds the monitoring workload according to a given number of nodes and metrics
- Scatters the monitoring tasks evenly across the available CPU cores as a multi-threaded code and Gathers the responses
- Each thread Queries and Collects monitoring data from monitored entities
- Interleaves BMC metrics (hardware data) with Resource Manager metrics (software data), and Writes metrics to InfluxDB

Data Processing

- Interleaves BMC metrics (hardware data) with Resource Manager metrics (software data), and Writes metrics to InfluxDB

At what time interval should we interleave the software data with hardware data?

INTERLEAVING

- At the same time interval with one we fetch hardware data ? i.e. about 6 secs
  - Not necessary, the software data may not change as frequently as hardware data
  - May cause extra overhead on the resource manager(UGE/SLURM)
- At a relatively large time interval, e.g. 1mins, 5mins
  - Not accurate, some jobs may start and finish during this large time interval
  - The interleaved metrics are affected by how we choose the time interval

INTERLEAVING

- Interleaving software data with hardware data as needed
- Retrieve data from log file directly instead of from resource manager database
  - Includes resource snapshots of running jobs and other cluster statics
  - All information is time related, whenever there's change, such as a new job submission or a job is finished, these kinds of info are logged into the reporting file

- Reporting file (UGE)[1]
- Job accounting log file (SLURM)[2]

Ref: [1] http://www.gridengine.eu/mangridengine/htmlman5/reporting.html, [2] https://www.mankier.com/1/sacct

WORKFLOW

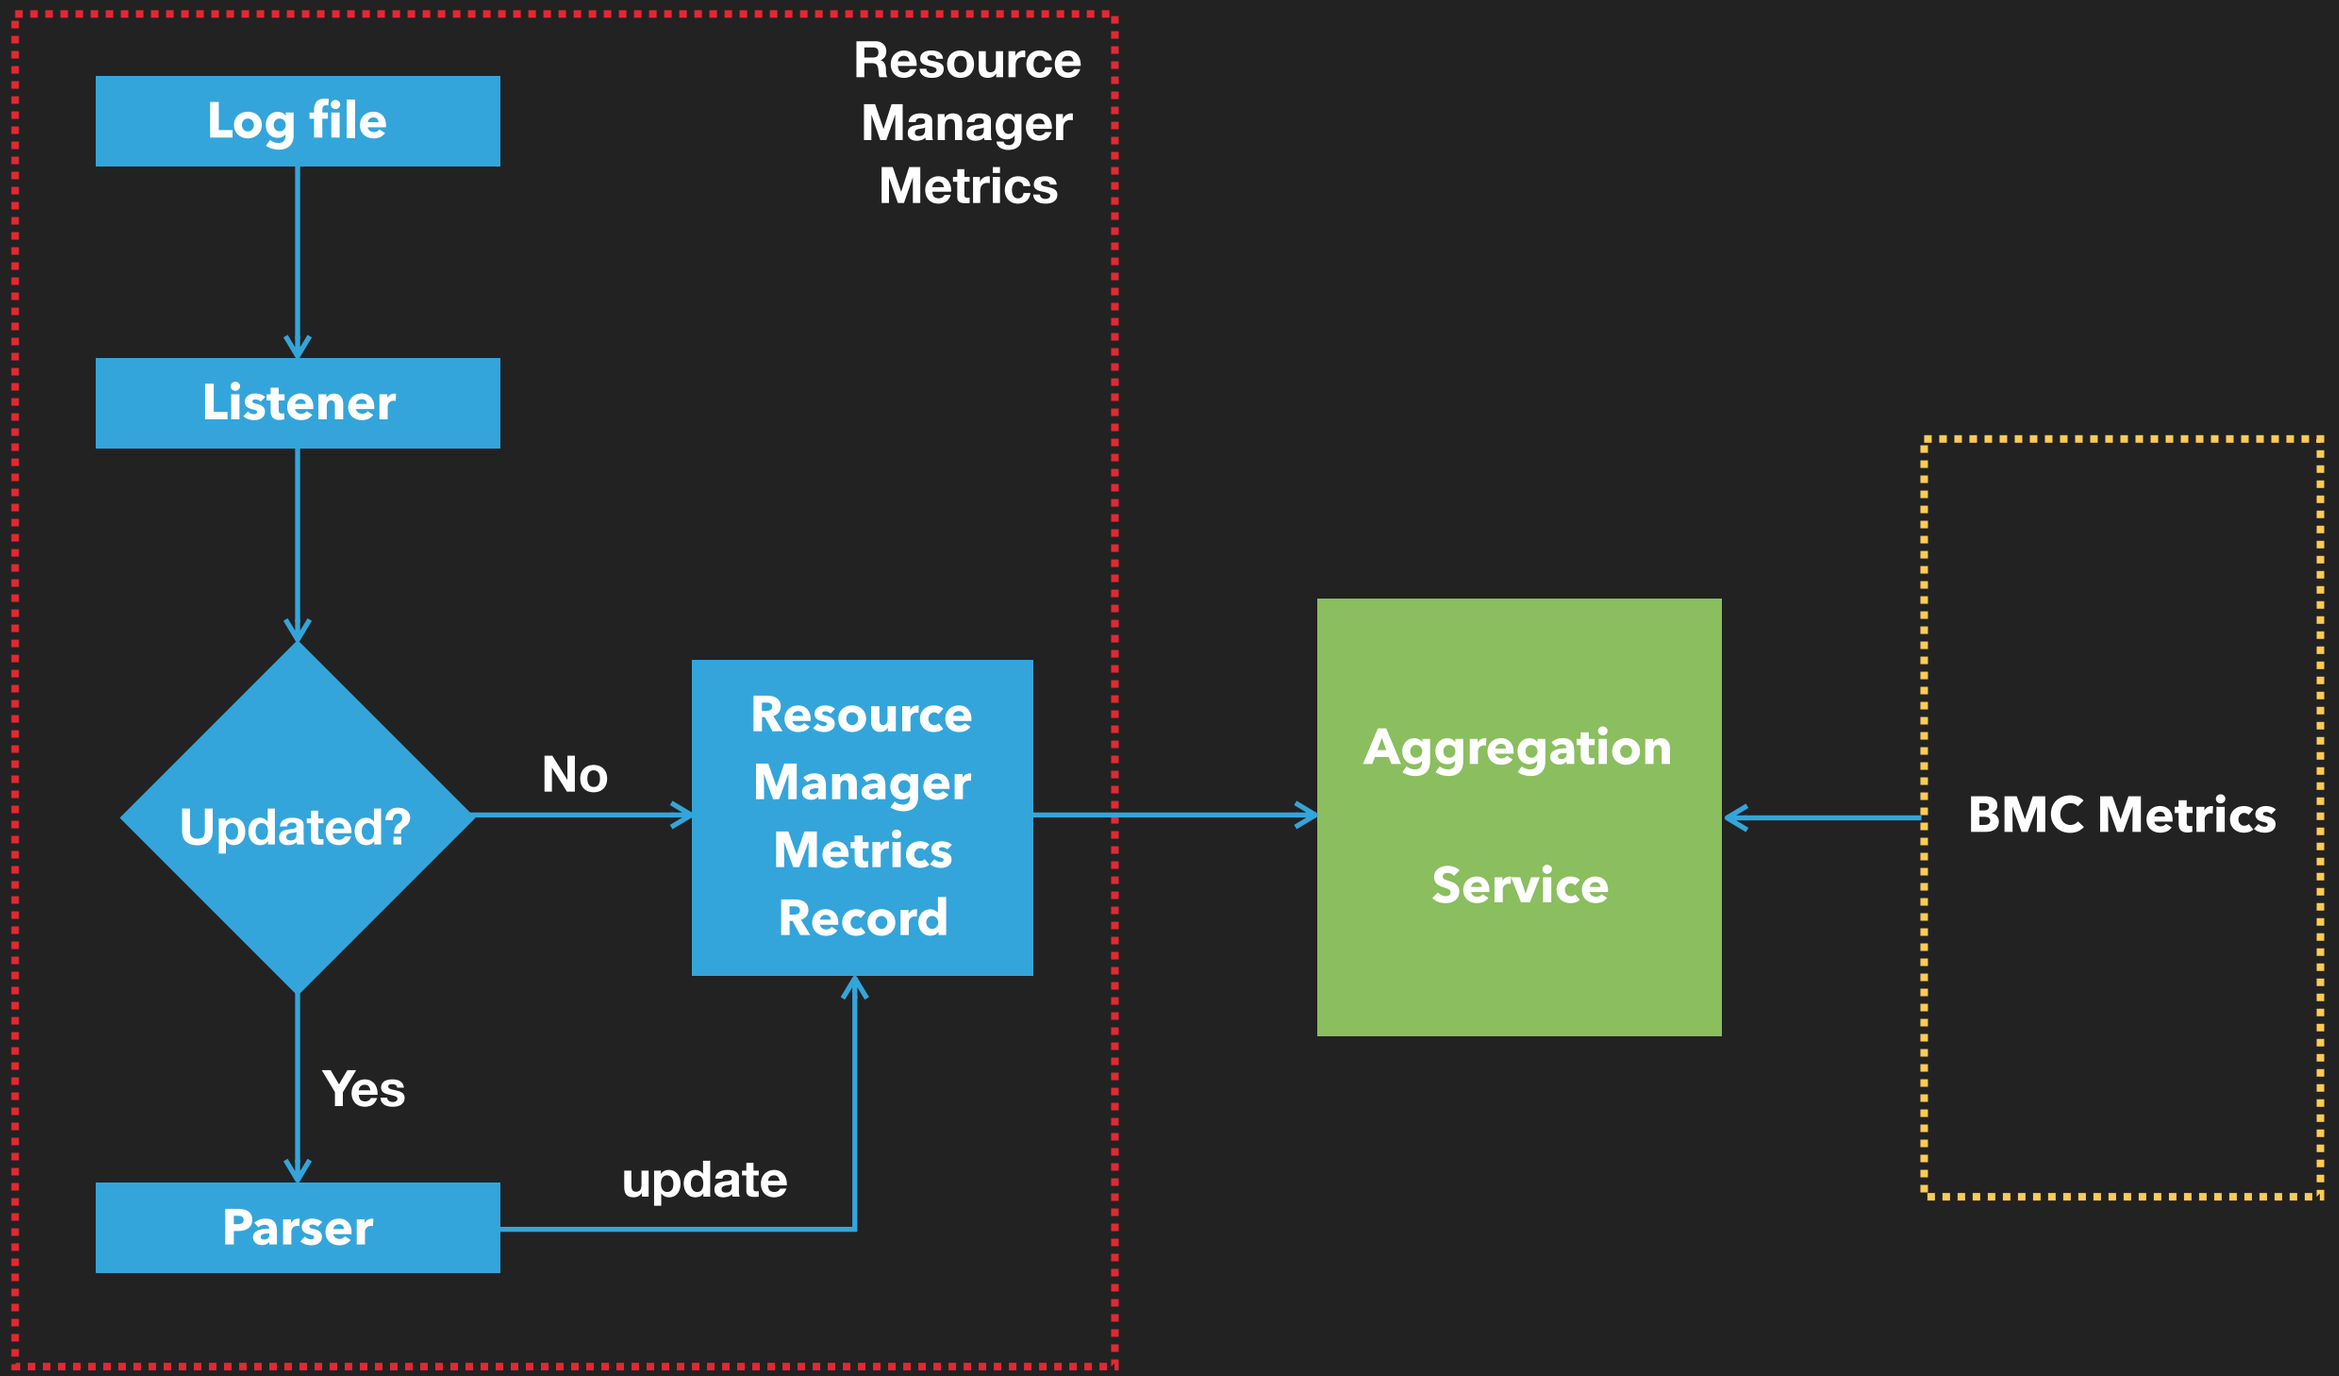
Resource Manager Metrics
Log file
Listener
Updated?
No
Resource Manager Metrics Record
Yes
Parser
update
Aggregation Service
BMC Metrics

AGGREGATION

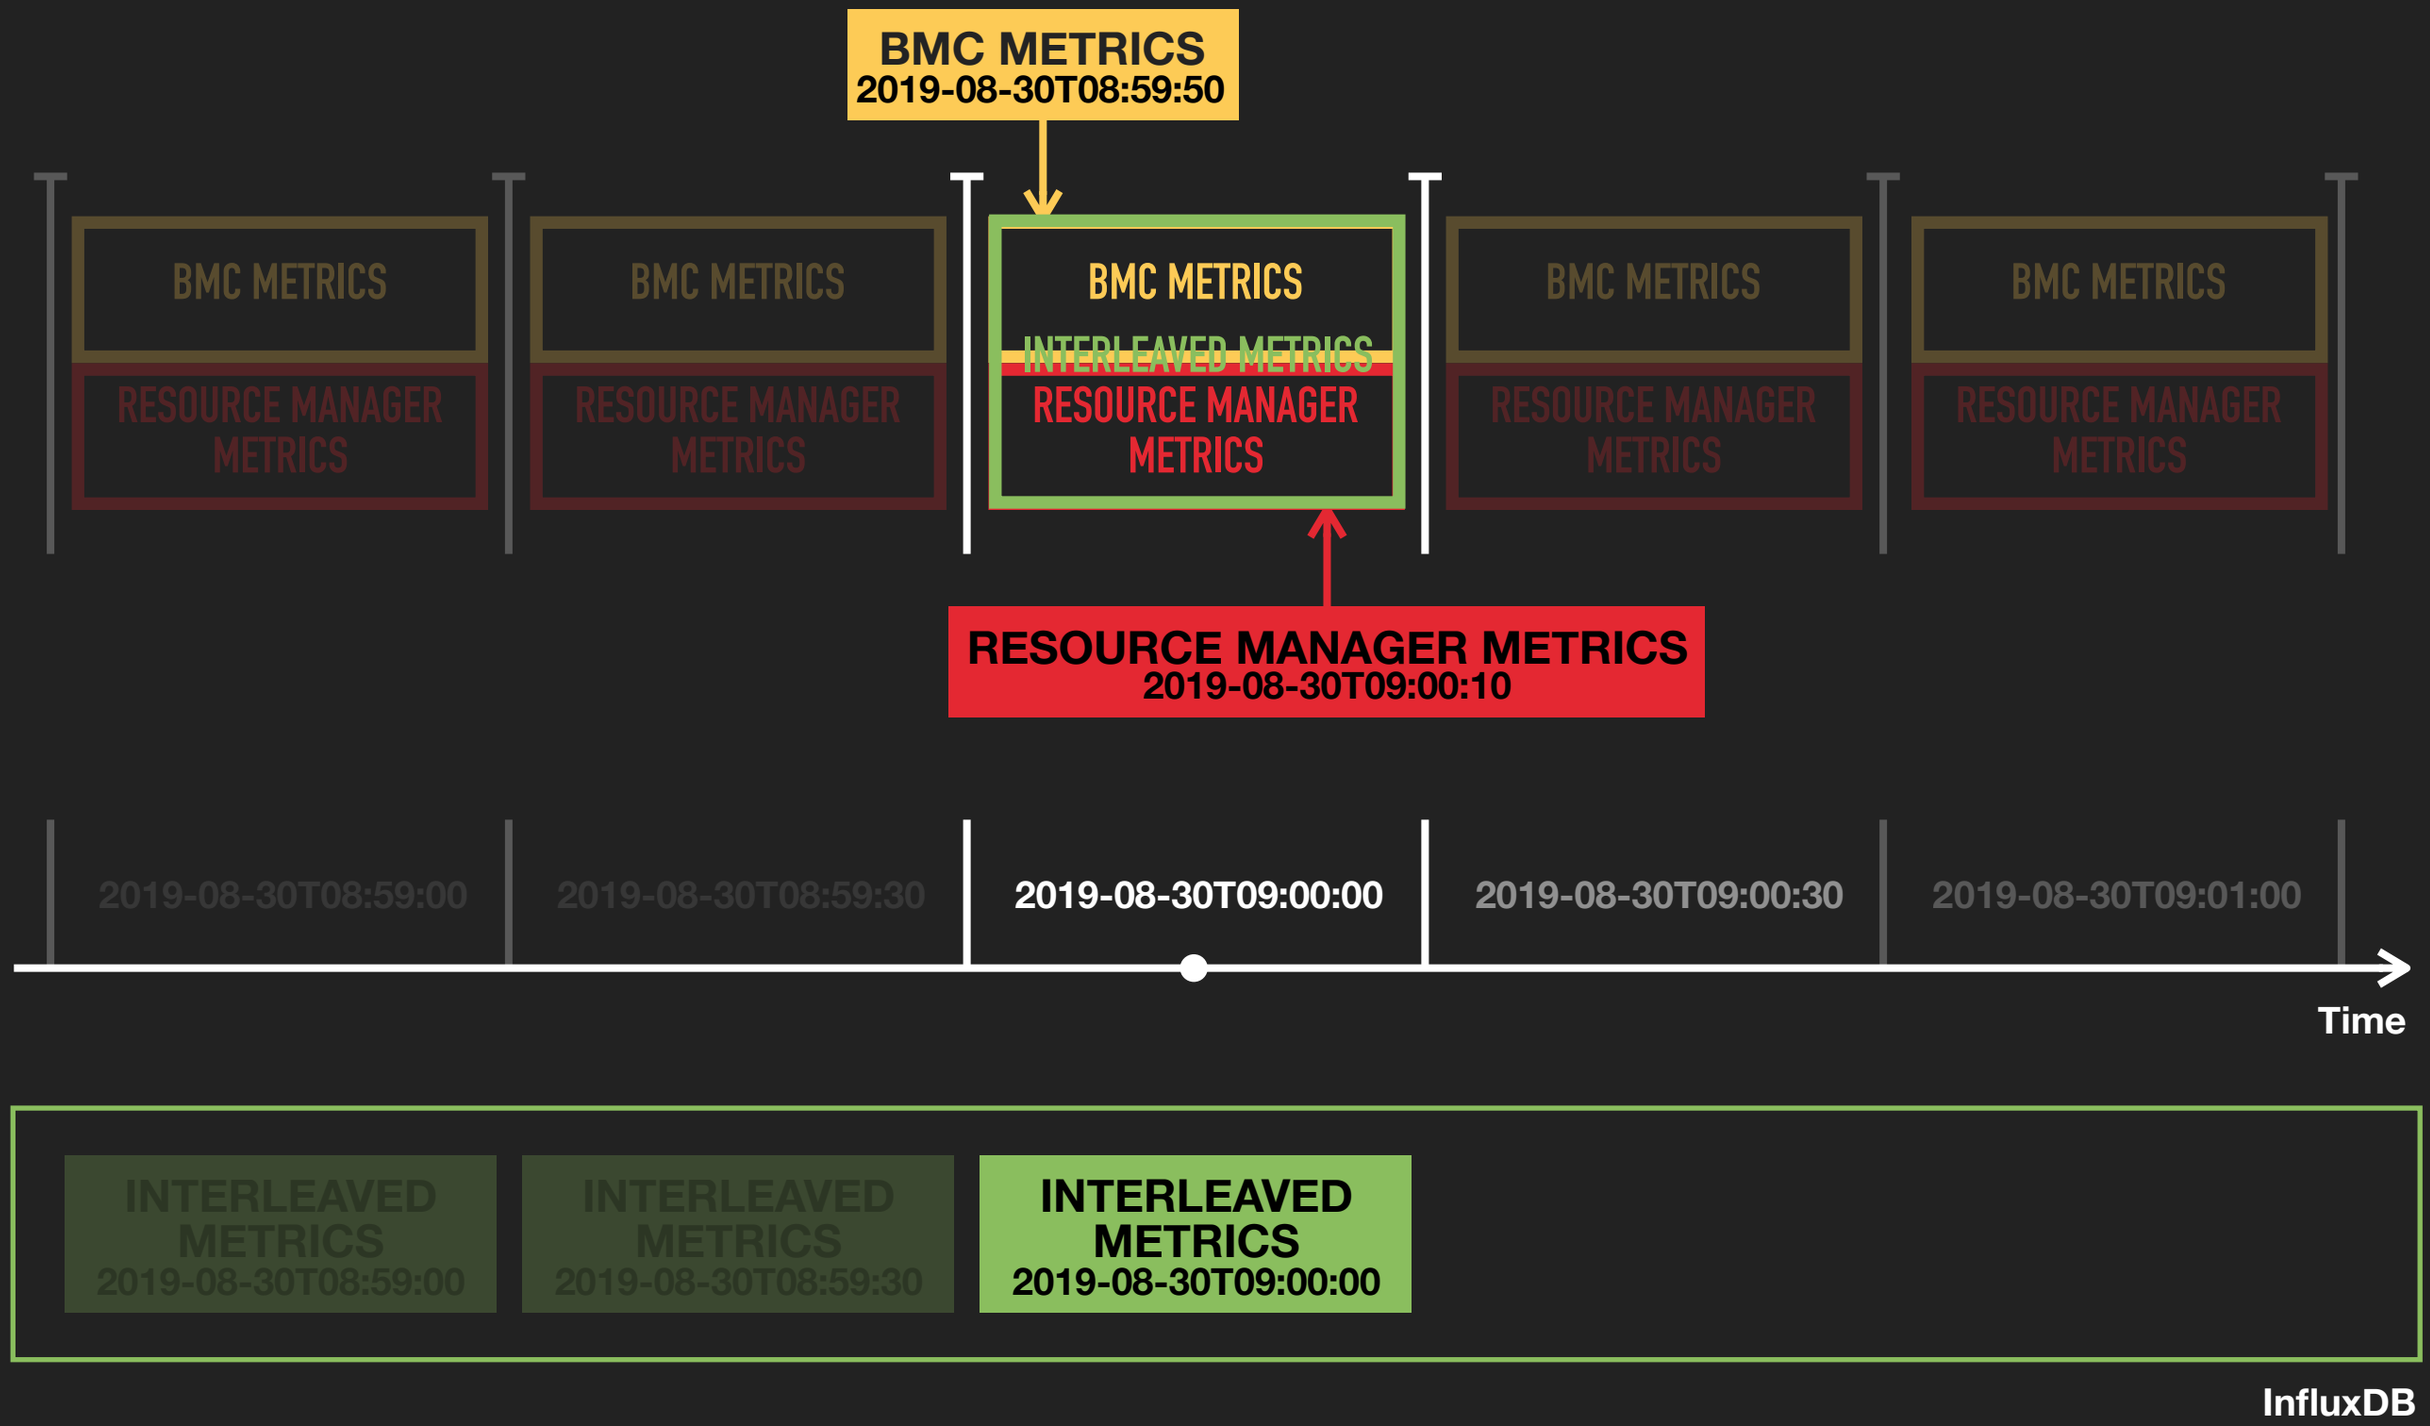

QUESTIONS?/COMMENTS?