

NSF/IUCRC CAC PROJECT

INTEGRATED VISUALIZING, MONITORING, AND MANAGING HPC SYSTEMS

Jie Li

Doctoral Student, TTU

11/20/2020

Advisors:

Mr. Jon Hass, SW Architect, Dell Inc.

Dr. Alan Sill, Managing Director, HPCC, TTU

Dr. Yong Chen, Associate Professor, CS Dept, TTU

Dr. Tommy Dang, Assistant Professor, CS Dept, TTU


INFLUXDB VS TIMESCALE

- ▶ InfluxDB vs TimescaleDB
 - ▶ Data model
 - ▶ Stability
 - ▶ Performance

DATA MODEL

InfluxDB

Non-relational DB (built from scratch in Go)
floats, ints, strings, and booleans
Only tags values are indexed



Time	Tag - NodeIP	Field - JobList
1583792296	"101.10.1.1"	"[123456, 123457]"



TimescaleDB

Relational DB (built on PostgreSQL)
floats, ints, strings, booleans, arrays, JSON, etc.
Indexable on all fields



Time	Field - NodeIP	Field - JobList
1583792296	"101.10.1.1"	[123456, 123457]

NodeIP	Cluster	Rack	CPUs	...
"101.10.1.1"	"Quannah"	1	36	...

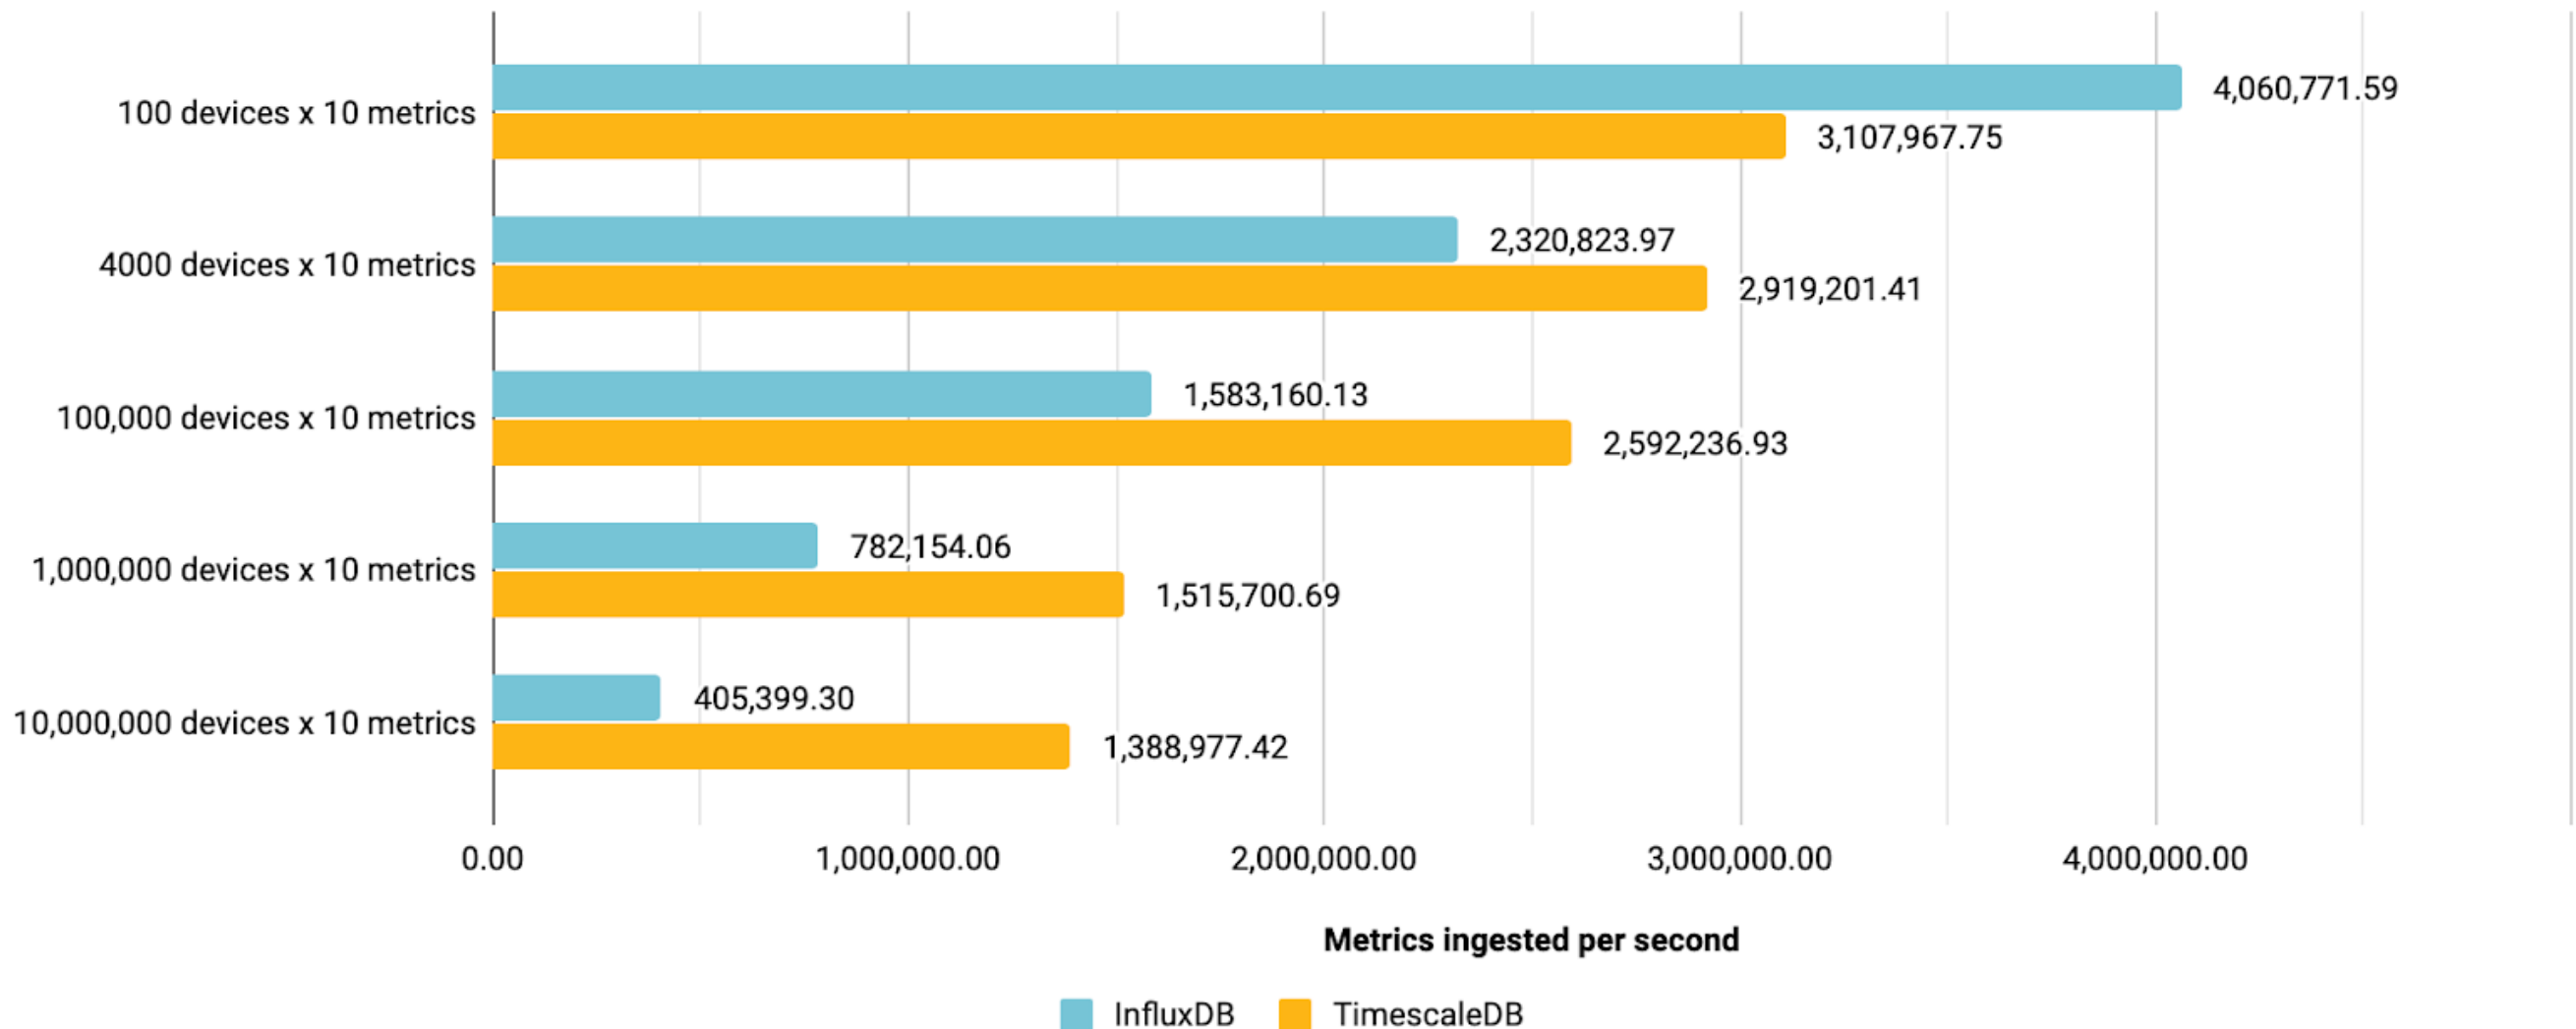
JobId	JobName	Start	NodeList	...
123456	"test"	1583792200	36	...

STABILITY

- ▶ Inserting batches into InfluxDB
 - ▶ Inserting batches of 10k into InfluxDB at high cardinalities will have **write errors** caused by timeouts, exceeding the maximum cache memory size, fatal out of memory errors.
 - ▶ Increasing maximum cache size and decreasing the batch size could solve these errors.
- ▶ Reading queries on InfluxDB
 - ▶ InfluxDB at high cardinalities could consume all available memory to run the query and **crashed with an Out of Memory error.**
- ▶ Writing large batches and reading queries on TimescaleDB do not have such issues. **PostgreSQL limits system memory usage** with settings like `shared_buffers` and `work_mem`.

```
1ea-a741-0242ac110002 2888077
[httpd] 206.81.15.50 - - [03/Aug/2020:16:41:40 +0000] "POST /write?consistency=all&db=benchmark HTTP/1.1" 204 0 "-" "tsbs_load_influx" 34d36805-d5a8-1
1ea-a74f-0242ac110002 2236706
[httpd] 206.81.15.50 - - [03/Aug/2020:16:41:40 +0000] "POST /write?consistency=all&db=benchmark HTTP/1.1" 204 0 "-" "tsbs_load_influx" 3493bafd-d5a8-1
1ea-a747-0242ac110002 2705985
[httpd] 206.81.15.50 - - [03/Aug/2020:16:41:40 +0000] "POST /write?consistency=all&db=benchmark HTTP/1.1" 204 0 "-" "tsbs_load_influx" 34d2fbf3-d5a8-1
1ea-a74e-0242ac110002 2359814
fatal error: runtime: out of memory
```

Ingest Rate Comparison: InfluxDB vs TimescaleDB




- ▶ InfluxDB outperforms TimescaleDB for workloads with low cardinality
- ▶ InfluxDB insert performance drops off dramatically as cardinality increases.
- ▶ TimescaleDB has ~3.5x the insert performance as InfluxDB

PERFORMANCE

Query Performance (measured in milliseconds)

	100 devices x 1 metric			100 devices x 10 metrics			4,000 devices x 10 metrics		
Simple rollups ¹	Influx	Timescale	Influx / Timescale	Influx	Timescale	Influx / Timescale	Influx	Timescale	Influx / Timescale
single-groupby-1-1-1	11.33	12.11	94%	5.49	7.76	71%	6.15	6.02	102%
single-groupby-1-1-12	32.87	13.36	246%	26.48	14.62	181%	32.61	22.68	144%
single-groupby-1-8-1	43.56	7.29	598%	13.04	10.17	128%	16.09	17.06	94%
single-groupby-5-1-1	—	—	—	12.4	6.67	186%	14.76	8.62	171%
single-groupby-5-1-12	—	—	—	82.8	17.87	463%	106.8	23.08	463%
single-groupby-5-8-1	—	—	—	49.32	12.51	394%	64.6	17.53	369%
Aggregates ²									
cpu-max-all-1	—	—	—	13.84	13.69	101%	16.14	17.68	91%
cpu-max-all-8	—	—	—	95.36	56.61	168%	104.25	66.79	156%
Double rollups ³									
double-groupby-1	500.55	272.46	184%	152.64	331.54	46%	6,050.85	11,060.68	55%
double-groupby-5	—	—	—	703.36	508.7	138%	31,801.62	22,479.91	141%
double-groupby-all	—	—	—	1393.91	869.81	160%	65,212.69	34,603.17	188%
Thresholds ⁴									
high-cpu-1	2,652.17	304.9	870%	2,952.15	836.49	353%	180,235.94	35,049.85	514%
high-cpu-all	20.68	8.25	251%	29.5	11.42	258%	30.56	17.44	175%
Complex queries									
lastpoint ⁵	367.45	7.55	4,867%	192.69	9.49	2,030%	10,514.64	147.36	7,135%
groupby-orderby-limit ⁶	3,344.5	752.68	444%	2411.74	700.02	345%	114,419.32	27,990.85	409%

- ▶ Generally, Timescale outperforms InfluxDB.
- ▶ When simply rolling up a single metric, InfluxDB can sometimes outperform TimescaleDB
- ▶ TimescaleDB vastly outperforms InfluxDB for complex queries.

 InfluxDB outperforms
 TimescaleDB outperforms

CONCLUSION

- ▶ We **do NOT need to use a non-TSDB** to store static data (job data) if using TimescaleDB to store the HPC monitoring data.
- ▶ TimescaleDB is much more **stable** when writing and reading **high-cardinality** datasets.
- ▶ TimescaleDB **performs better** on writing and reading high-cardinality datasets.

We may use **TimescaleDB** as the main storage solution for monitoring the RedRaider cluster.



QUESTIONS?/COMMENTS?